

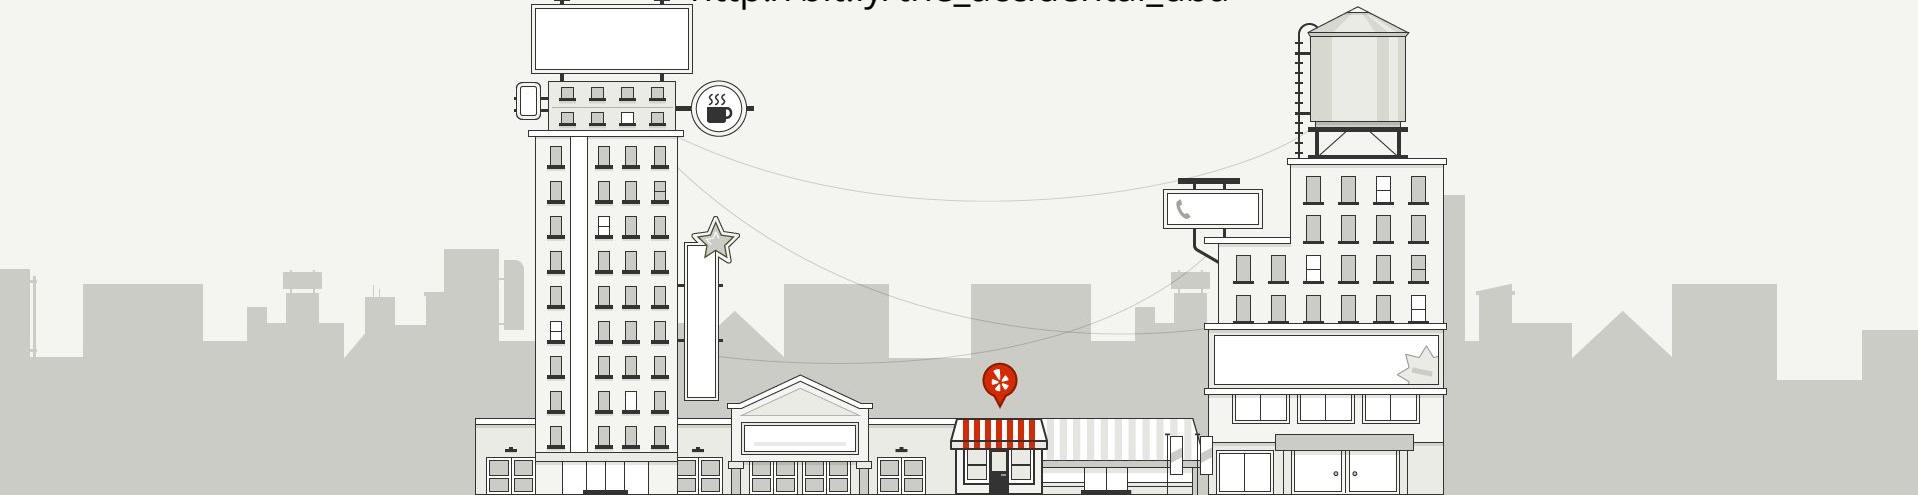
The Accidental DBA

Jenni Snyder

Yelp

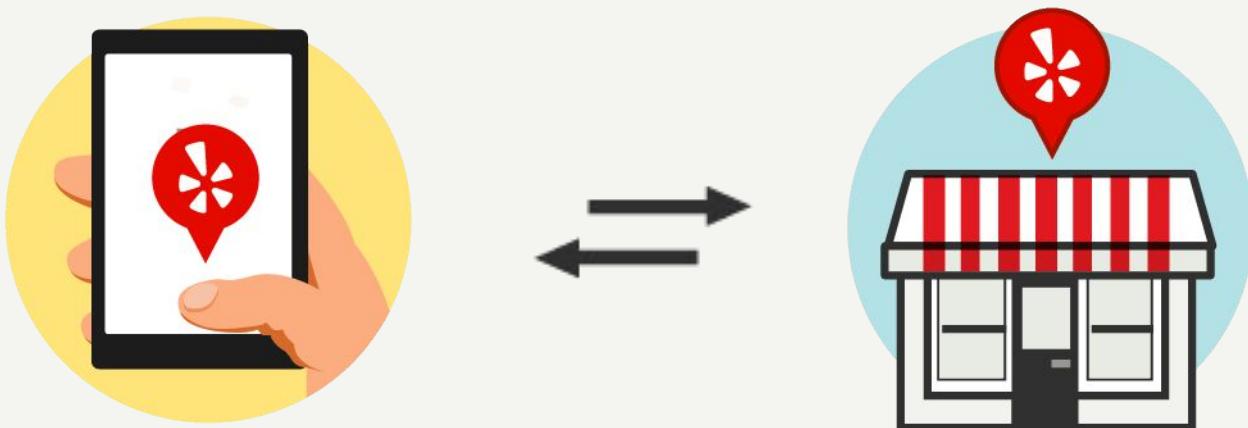
@jcsuperstar

http://bit.ly/the_accidental_dba



Yelp's Mission

Connecting people with great
local businesses.



Administrative

9:30am: Welcome & Let's talk about the basics

11:00am: Break

11:15am: Running MySQL in the real world

12:15pm: Final wrap up

12:30pm: Lunch



Agenda in More Detail

Part 1:

0. Brief Introduction
1. MySQL History, Philosophy, Community
2. Installation, Basic Configuration, Launching
3. Using it, Backing it up
4. Replication basics, Monitoring
5. Database Defense & fast feedback to developers



Agenda in More Detail

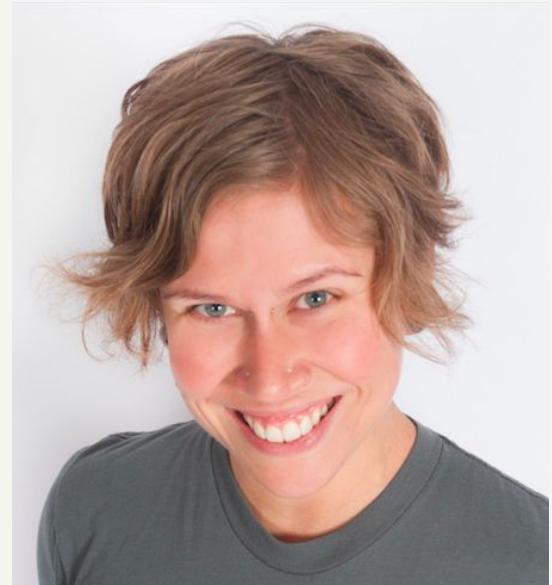
Part 2:

- Running in Production #realtalk
- Being Prepared
- Advanced Monitoring and Investigation
- Fantastic Errors and Where to Find Them
- When Things Are "Slow"
- Advanced Tuning
- Replication and Scaling Out
- Discussion

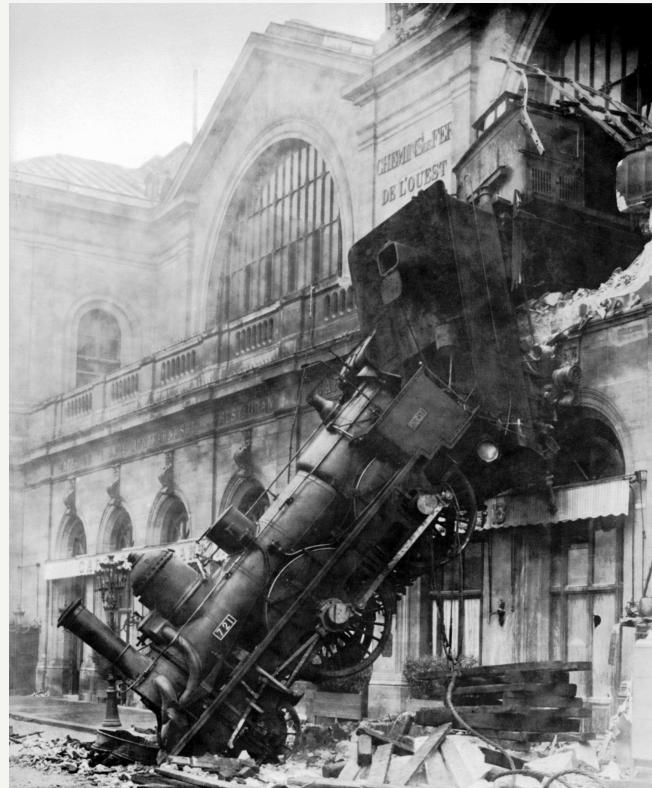


Who am I Why am I Here?

- My name is Jenni
- Engineering Manager @ Yelp
- First DBA there in 2011



Got Into Databases by Accident







Myth 1: *databases are scary!*

A photograph of a person from behind, wearing a dark green or black jacket, standing in a dense forest. The background is filled with out-of-focus tree branches and leaves.

Myth 2: DBAs are scary!





Why Databases?



It Was on Fire When I Got Here



You may already have a database.



And it's Not Clear Who Fixes it

"If someone is asking you who owns the database, it's probably you"

- @mipsytippsy



Don't Feel Bad!

- Yelp didn't have a DBA for >5 years
- We've acquired companies with great big DBs
- There are probably some easy things to fix

iii
nowait

turnstyle





White Chocolate Chip Cookies
£1.80 each
£8.00 per 10

Oat and Raisin Cookies
£1.80 each
£8.00 per 10

Milk Chocolate Chip Cookies
£1.80 each
£8.00 per 10

White Chocolate and Cranberry Cookies
£1.80 each
£8.00 per 10

Dark Chocolate Chip Cookies
£1.80 each
£8.00 per 10

Chocolate Chip Cookies
£1.80 each
£8.00 per 10

Just in Case You Don't Have a DB Yet...

Database Flavors:

- SQL
- NoSQL
- Geospatial
- Graph



SQL/Relational Databases

- MySQL
- PostgreSQL
- SQLite



NoSQL

- Cassandra
- MongoDB
- Couchbase
- Redis



Geospatial & Graph

- PostGIS
- Neo4j
- Titan



When Not to Use a Database

When it's a:

- log
- queue
- cache
- photo
- disk directory
- spreadsheet
- data warehouse/cube/shed/lake/garage/whatever

Use the right tool for the job.





Brief History of MySQL

- [MySQL - Wikipedia](#)
- First released in 1995
- Pluggable storage engines
- Written in the age of spinning rust
- Transactional engine **InnoDB** released in 2001
- Oracle buys InnoDB in 2005
- Bought by Sun in 2008
- Oracle bought Sun in 2010



MySQL - Why Does it Do What It Does

- It is meant to work best the most of the time
- It is ~~hard~~ impossible to optimize for all cases
- That sometimes surprises you

"Performs extremely well in the average case" - Wikipedia



The MySQL Community

- The manual is awesome
- [DBA StackExchange](#)
- So many forums - "just <search engine> for it"
- Conferences
- Lots of great open source



Highly Configurable

- [MySQL Manual - Server System Variables](#) - 500+
- 400+ status metrics

Don't freak out. You don't have to care about these yet :)





Installation

- Source, package, #whatever
- Drop a config in /etc/my.cnf:
 - basedir, datadir
 - socket
 - pid-file
 - **server-id**



MySQL Files - Logs

- Error log
- Slow log
- Don't log warnings for now



MySQL Files - Replication Logs

- You may have lots of these:
 - Binary logs
 - Relay logs



MySQL Config - memory

- `innodb_buffer_pool_size` ~ 75% memory
- Worry about the rest of it later





Let's Fire This Puppy Up

- `service mysql start`

or

- `/etc/init.d/mysql start`
- Runs on port 3306
- Throw a party
- Or just connect to it
 - Use the `mysql` command



```
jjsnyder@db:~$ ps auxw| fgrep mysql
```

root	16661	0.0	0.0	4460	1692	?	Ss	Oct09	0:00
	/bin/sh	/usr/bin/mysqld_safe							
mysql	17877	1.2	88.8	30221480	27893084	?	S1	Oct09	318:00
	/usr/sbin/mysqld	--basedir=/usr							
	--datadir=/nail/databases/mysql/data								
	--plugin-dir=/usr/lib/mysql/plugin	--user=mysql							
	--log-error=/nail/databases/mysql/log/err.log								
	--pid-file=/nail/databases/mysql/mysqld.pid								
	--socket=/nail/databases/mysql/mysqld.sock	--port=3306							

```
jsnyder@db:~$ mysql -u root
```

Welcome to the MySQL monitor. Commands end with ; or \g.

Your MySQL connection id is 1124438

Server version: 5.6.32-78.1-log Percona Server (GPL), Yelp Release 1

Copyright (c) 2009-2017 Percona LLC and/or its affiliates

Copyright (c) 2000, 2017, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

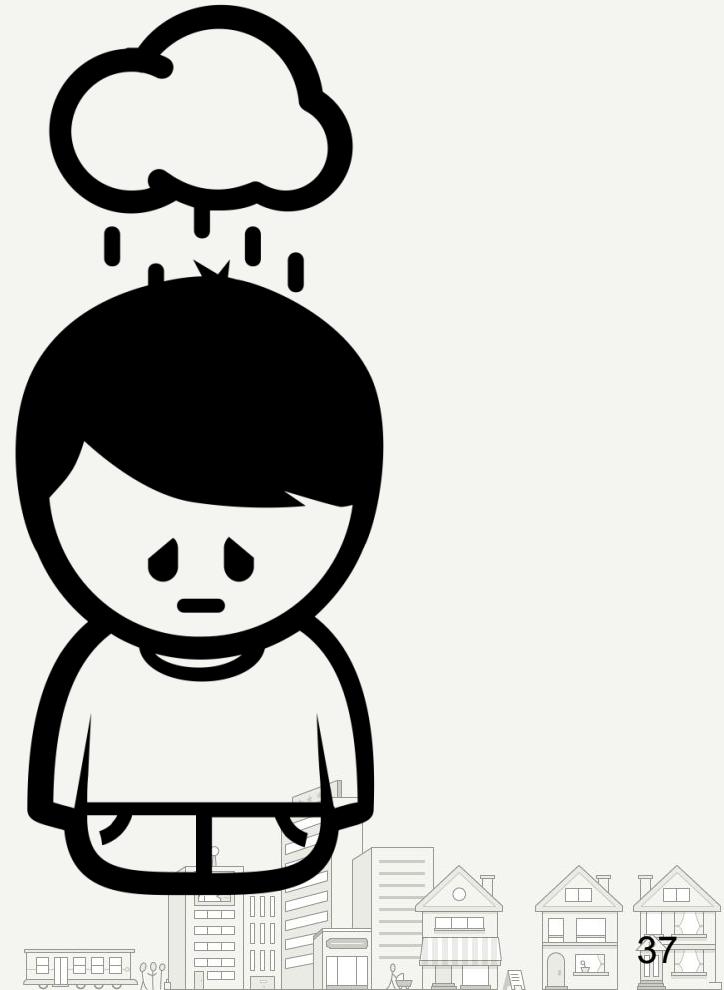
Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql>
```

```
mysql> show databases ;
+-----+
| Database      |
+-----+
| information_schema |
| mysql          |
| performance_schema |
| sys            |
+-----+
4 rows in set (0.02 sec)
```

When it Won't Start

- Find the start script
- Find that err.log
- Run `mysqld_safe` manually





MySQL Grants

- MySQL uses GRANTS
- Users are a combination of:
 - User (name)
 - Password
 - Host - % is wildcard
- localhost uses unix domain socket
- 127.0.0.1 uses TCP/IP



MySQL Grant Levels

- GRANTS can be for different levels:
 - Global `*.*`
 - Database `yelp_db.*`
 - Table `yelp_db.user`



```
mysql> CREATE USER 'yolo'@'%' IDENTIFIED BY '' ;
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> GRANT ALL ON *.* TO 'yolo'@'%' WITH GRANT OPTION ;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON first_db.* TO
'app'@'10.10.10.%' IDENTIFIED by 'sweetpassword' ;
Query OK, 0 rows affected (0.00 sec)
```

```
mysql> GRANT CREATE, DROP, INDEX, TRIGGER, SELECT, INSERT,
UPDATE, DELETE ON *.* TO 'schema_manager'@'10.%' IDENTIFIED BY
'yougettheidea' ;
Query OK, 0 rows affected (0.01 sec)
```

```
mysql> select user, host from mysql.user ;
+-----+-----+
| user          | host        |
+-----+-----+
| yolo          | %           |
| schema_manager | 10.%       |
| app           | 10.10.10.% |
| mysql.sys     | localhost   |
| root          | localhost   |
+-----+-----+
5 rows in set (0.00 sec)
```

```
mysql> show grants for 'root'@'%';
+-----+
| Grants for root@% |
+-----+
| GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' WITH GRANT OPTION |
| GRANT PROXY ON ''@'' TO 'root'@'%' WITH GRANT OPTION |
+-----+
2 rows in set (0.00 sec)
```

```
mysql> show grants for current_user ;
+-----+
| Grants for root@% |
+-----+
| GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' WITH GRANT OPTION |
| GRANT PROXY ON ''@'' TO 'root'@'%' WITH GRANT OPTION |
+-----+
2 rows in set (0.00 sec)
```

Debugging Connection Issues

- Make sure MySQL is running
- On that host & port
- Match user & host in error with grants
- Read err.log
- Use skip-grant-tables to start over



What if You Don't Know the Password?

- Add skip-grant-tables to your my.cnf
- Restart MySQL

ALTER USER 'root'@'localhost' IDENTIFIED BY 'MyNewPass';



Debugging MySQL Connection Errors

- Start with broad grants ('user'@'%')
- Refine from there



```
mysql -u myuser -p mypass -h myhost -P myport
```

```
jsnyder@db:~$ mysql -u myuser -h myhost  
ERROR 1045 (28000): Access denied for user myuser@'10.10.10.10'  
(using password: NO)
```

```
jsnyder@db:~$ mysql -u myuser -p -h myhost  
Enter password:  
ERROR 1045 (28000): Access denied for user  
'myuser'@'10.10.10.10' (using password: YES)
```

```
jsnyder@db:~$ mysql -u myuser -p -h myhost  
Enter password:  
ERROR 2005 (HY000): Unknown MySQL server host 'myhost' (0)
```



Developing Against MySQL

We use:

Language	Python
MySQL Client Library	mysqlclient-python (there is also Connector/Python)
Object Relational Mapping (ORM)	SQLAlchemy
Schema Manager	Hand-rolled system

To ORM or not to ORM?

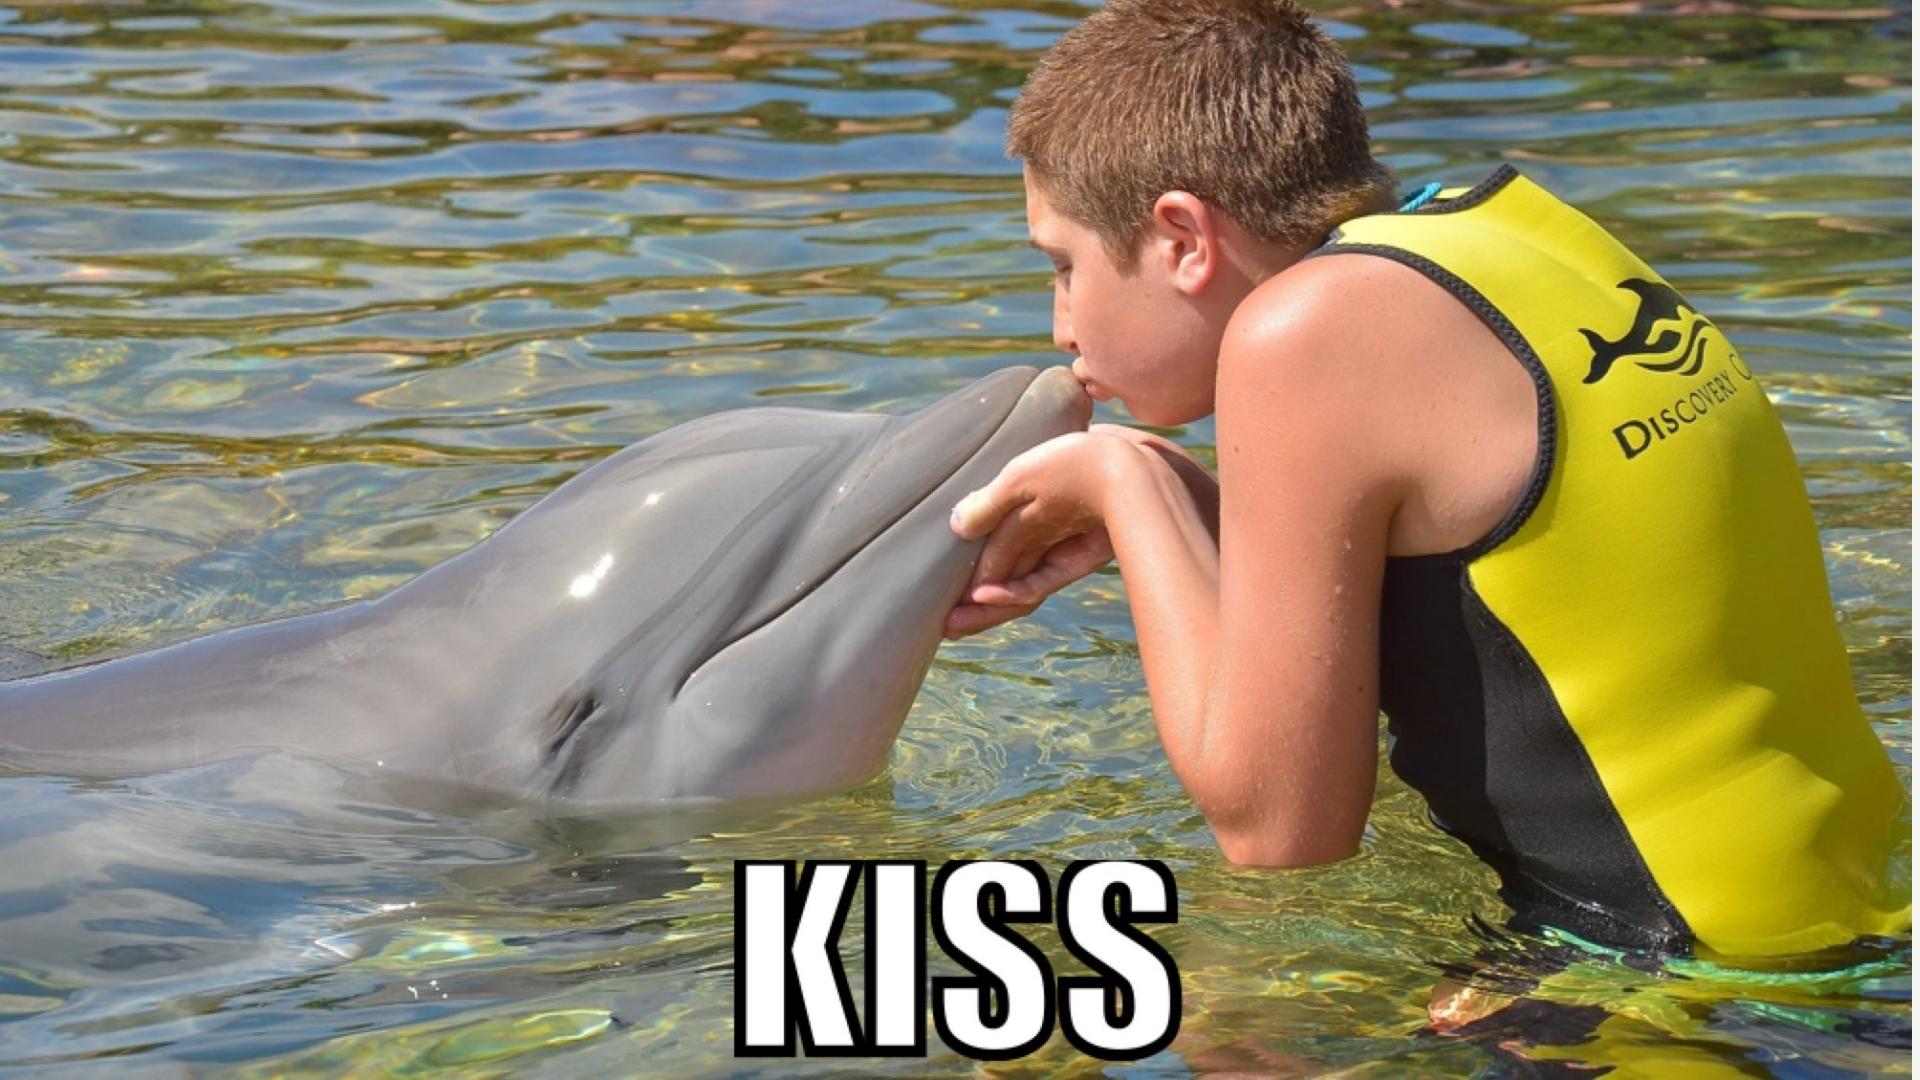


About That Schema

- Have developers read up on normal form
 - It's OK to denormalize for performance (later)
- Use one table for one thing
- Try not to get too fancy right away

Reference: [Database Normalization Basics](#)





KISS

When Developing Against MySQL

- Don't swallow database client warnings/errors
- Track client error rates
- Make sure devs get error feedback

Build feedback to go directly to developers.



Application Connections

- You can start simple
- Reads and writes scale differently
- Use discovery, load balancer or proxy
 - Smartstack
 - [ProxySQL](#)
 - [MySQL Router](#)

Use technology you already have if possible.





Back to the Administration Part

- When you care about your database
 - Back it up to a safe place
 - Test restoring them
 - Have more than one



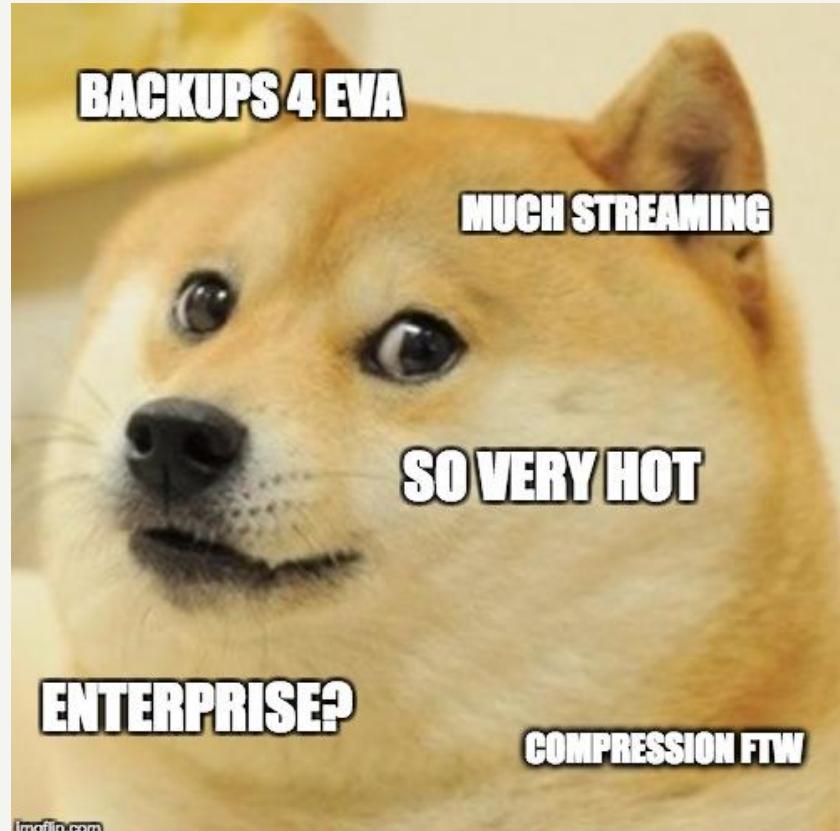
Different Kinds of Backups

- Binary a.k.a. hot backups
- Logical backups
- Point in time recovery
 - tl;dr: backup your binary logs



Binary Backups

- "They're Hot"
- MySQL Enterprise
- [xtrabackup](#)
 - Streaming



Logical Backups

- Spoiler alert: it's SQL
- `mysqldump`
- Best run on an offline/not replicating DB
- Sometimes the only way to upgrade



Backing up the Binary Logs

- Logs of all changes
- Can be used to "roll" a DB forward in time
- More on these later...
- Copy & compress

Sample command to look at these:

```
mysqlbinlog --base64-output=DECODE-ROWS -v mysql-bin.000001
```



Testing Backups

- Backups are only as good as your ability to use them
- Test your restores



Point in Time Recovery With Binary Logs

- `mysqlbinlog` turns binary logs into SQL
- Use the coordinates you saved
- It's [documented](#), but I'm not aware of a great tool



```
j snyder@db:~$ cat  
/path/to/backup/xtrabackup_binlog_info  
mysql-bin.000003      589767
```

```
j snyder@db:~$ mysqlbinlog  
/path/to/binlogs/mysql-bin.000003  
/path/to/binlogs/mysql-bin.000004 \  
--start-position=589767 \  
--stop-datetime="2017-10-29 09:00:00" | mysql  
<args>
```

Practice!

- Try it before you need it
- We call these "war games"
 - And have failover exercises
- The best plan is one you know will work!

More on replication next.





But First, a Word About Words

MySQL uses the terms *MASTER* and *SLAVE*.

It is modern and inclusive to use *PRIMARY* and *REPLICA*.

This talk uses MASTER and REPLICA.



Replication

- Basics
 - One master + multiple replicas
 - Client-initiated, client configured
 - The replication user also needs grants

```
GRANT REPLICATION CLIENT, REPLICATION SLAVE  
ON *.* TO 'replication'@'10.%' ...
```



Logs & Threads

- Master executes statement, writes it to:
 - Binary log
- Replica's IO_THREAD reads from master, writes it to:
 - Relay log
- Replica's SQL_THREAD executes statement, writes it to:
 - Binary log

Replicas can have their own replicas!



Basic Configuration

- Unique server-id is a *must*
- Binary log formats
 - STATEMENT
 - ROW
 - MIXED
- Log-slave-updates = ON
- master_info_repository = TABLE

MySQL won't execute events from it's own server-id.



```
jsnyder@db:~$ fgrep bin /etc/my.cnf
# binary logging is required for replication
log-bin = ./master_log/mysql-bin
relay-log = ./relay_log/relay-bin
binlog_format = statement
```

```
jsnyder@db:~$ fgrep server-id /etc/my.cnf
server-id = 171739394
```

```
jsnyder@db:~$ fgrep slave /etc/my.cnf
log_slow_slave_statements = 1
log-slave-updates = 1
slave_net_timeout = 30
slave-skip-errors = 1141
```

Configuring Replication

- CHANGE MASTER TO
 - MASTER_USER =
 - , MASTER_PASSWORD =
 - , MASTER_HOST =
- If using GTID:
 - , MASTER_AUTO_POSITION = 1
- If not:
 - , MASTER_LOG_FILE =
 - , MASTER_LOG_POSITION =



Start it up

- START SLAVE
- SHOW SLAVE STATUS



```
mysql> SHOW SLAVE STATUS \G
```

```
***** 1. row *****  
Slave_IO_State: Waiting for master to send event  
Master_Host: 10.10.10.44  
Master_User: replication_user  
Master_Port: 3306  
Master_Log_File: mysql-bin.000008  
Read_Master_Log_Pos: 914832470  
    Relay_Log_File: relay-bin.000014  
    Relay_Log_Pos: 914832633  
Relay_Master_Log_File: mysql-bin.000008  
Slave_IO_Running: Yes  
Slave_SQL_Running: Yes
```

```
...
```

```
mysql> SHOW SLAVE STATUS \G
... (continued)

                Last_Error: 0
                Last_SQL_Error:
...
                Exec_Master_Log_Pos: 914832470
...
                Seconds_Behind_Master: 0
...
                Last_IO_Error: 0
                Last_IO_Error:
                Last_SQL_Error: 0
                Last_SQL_Error:
```

What Could Possibly go Wrong?

- Replication errors in two places
- `err.log`

```
2017-10-28 17:22:06 5061 [ERROR] Slave SQL: Error 'Duplicate entry  
'14' for key 'PRIMARY'' on query. Default database: 'yelp_db'.  
Query: 'INSERT INTO user (id, name) VALUES ('14', 'Barbara')',  
Error_code: 1062
```

- `SHOW SLAVE STATUS:`

`Last_SQL_Errno: 1062`

`Last_SQL_Error:'Duplicate entry '14' for key 'PRIMARY'' on query.
Default database: 'yelp_db'. Query: 'INSERT INTO user (id, name)
VALUES ('14', 'Barbara')'`



Try Breaking Errors Up

- Who had the problem?

[ERROR] Slave SQL: Error

- What was it trying to do?

Query: 'INSERT INTO user (id, name) VALUES ('14', 'Barbara')'

- What happened?

'Duplicate entry '14' for key 'PRIMARY'' on query.

- Where was this?

Default database: 'yelp_db'



Common Causes of Errors

- Wrong binary log coordinates
- Inconsistent data





Monitoring - Server

- If you care about it, monitor it
- Lots of open source options
 - Built-in to your monitoring solution
 - percona-toolkit
 - Cheap scripts
 - Look on github

Humans like ~~pictures~~ graphs!



Host Stuff

- Disk space:
 - We file Jira ticket at 85%
 - And page at 95%
- Is MySQL running, does it seem OK?
- Are its cron jobs completing?
- Does its config match running state?



MySQL Server Stuff

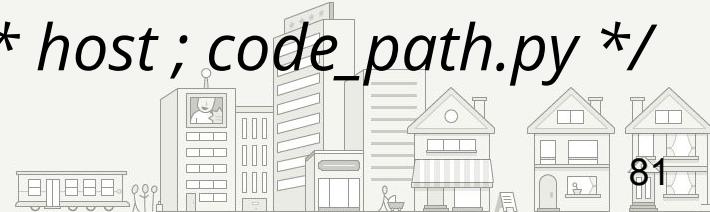
- Auto-increment id limits
- Binary log growth
- Killed queries
- Replication hierarchy
- `read_only`
- Temporary tables on disk
- Threads connected
- Threads running
- More fancy replication checks



MySQL Client Stuff

- Have your application:
 - Track client errors
 - Track query times
 - Bonus: query annotations

select id, name from user where id = 1 / host ; code_path.py */*





Automated Defense

Kill and Log:

- Long-running queries
- Long-running transactions
- Too many of the same queries

Send feedback directly to developers via the client.



Handling Client Errors

- MySQL client will receive errors
 - And retry
 - Or fail *gracefully*
- Track client errors & statistics
- Alert devs where reasonable

Let developers see immediate feedback of a change!



Automated Testing

Before devs push to production they can test for:

- Queries that will be slow
- Transactions that do too much

Send feedback directly to developers via the client.



Is All This Really Worth it?

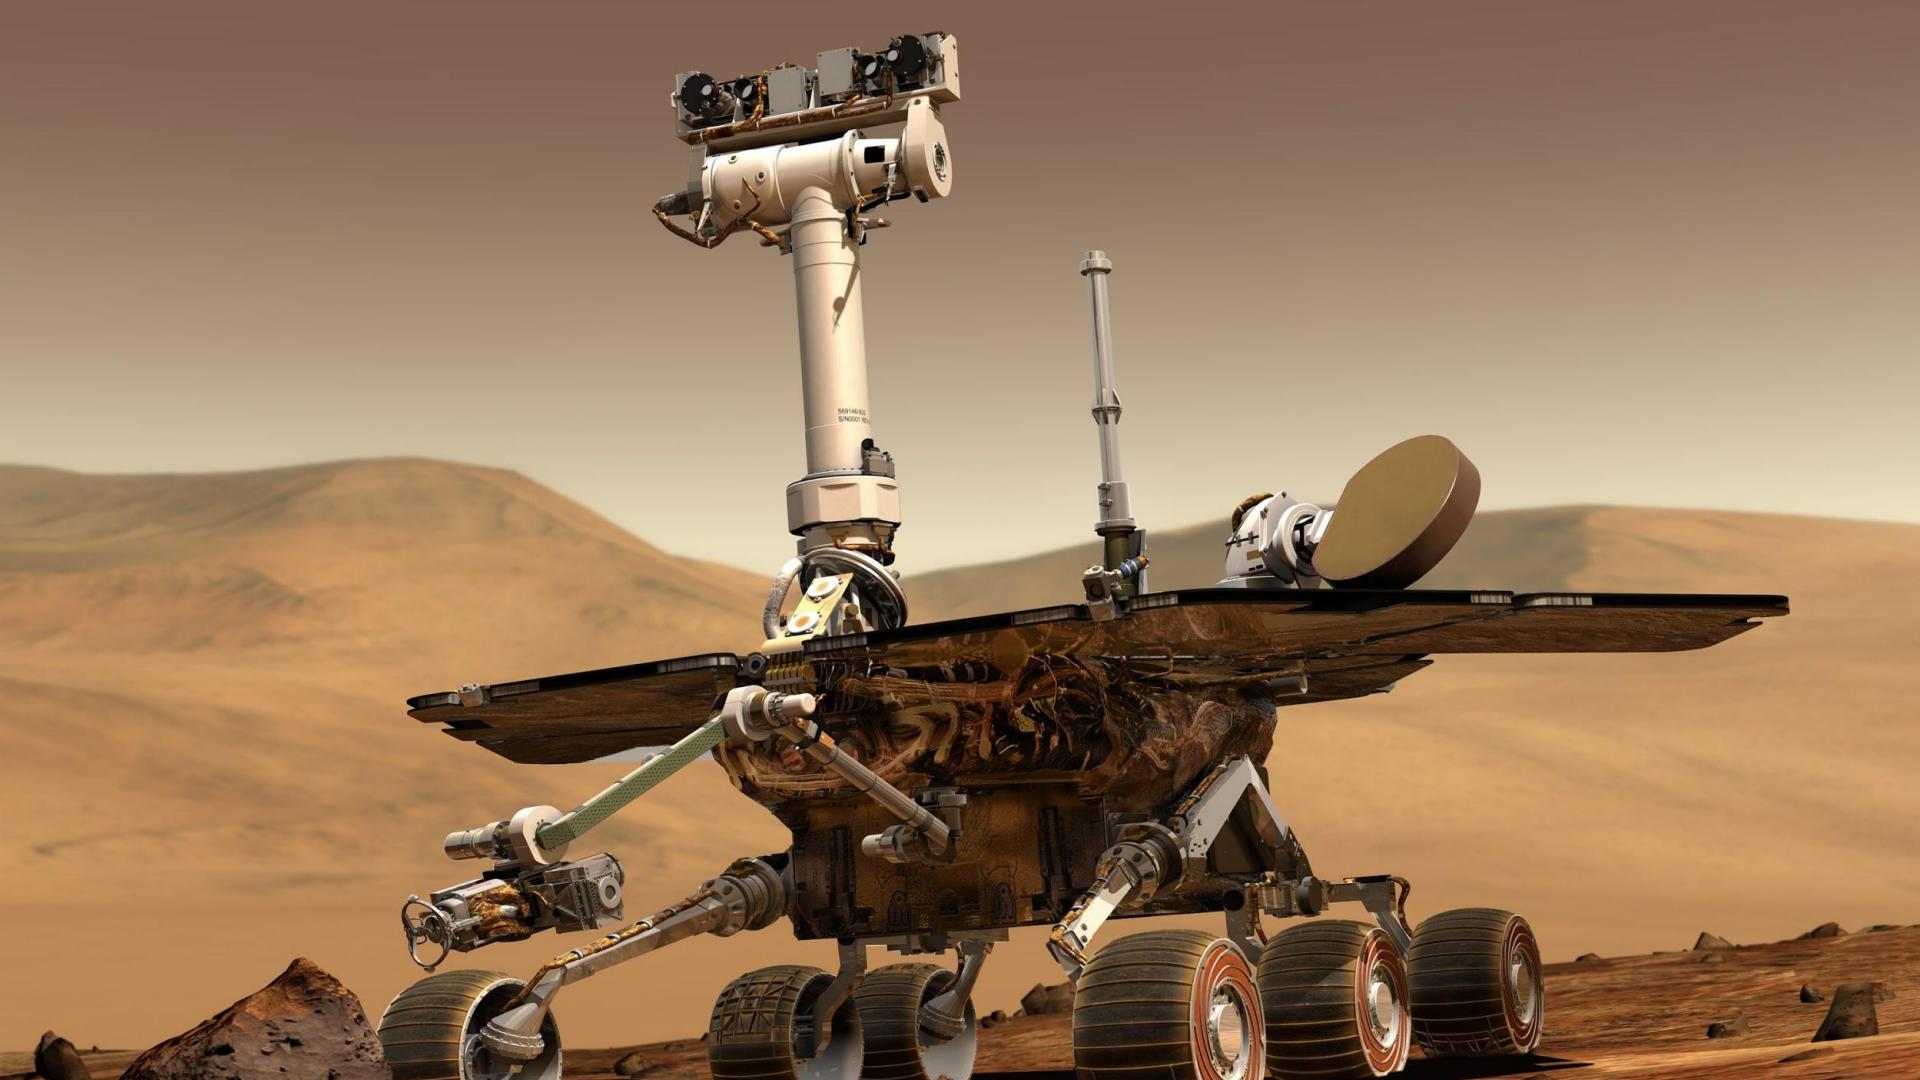
Seems like a lot of work!

Let's moment to reflect on humans and feelings,









Think About it - Ops Perspective

If you get paged:

- And it was because of someone else's change
- You might feel #negative_emotion
- Especially if it happens again...



Think About It - Dev Perspective

If someone tells you they got paged:

- You do something about it, you didn't know
- But if it happens again...
- You might feel #negative_emotion



Useful Automation/Monitoring

Will:

- Let developers discover issues & fix on their own
- Faster iteration while developing
- Confidence when pushing
- Keep the pager for when you're really needed

<3 <3 <3





**Questions
Answers**

Let's Talk About It

Questions, comments, quemments?

We covered:

1. MySQL History, Philosophy, Community
2. Installation, Basic Configuration, Launching
3. Using it, Backing it up
4. Replication basics, Monitoring
5. Database Defense & fast feedback to developers



Before we Break

Here is what we'll cover when we get back

- Running in Production
- Being Prepared
- Advanced Monitoring and Investigation
- Fantastic Errors and Where to Find Them
- When Things Are "Slow"
- Advanced Tuning
- Replication and Scaling Out
- Discussion





And we're back

Part 2a:

- Running in Production
- Being Prepared
- Advanced Monitoring and Investigation
- Fantastic Errors and Where to Find Them



And we're back

Part 2b:

- When Things Are "Slow"
- Advanced Tuning
- Replication and Scaling Out
- Discussion





Welcome to Production!

Some things to think about:

- Expectations
- SLOs - uptime and query time
- Caching where you can
- Degrading gracefully

Be prepared, or prepare to be surprised!



Because We Measure Uptime in 9's

Nothing is up 100%

- Get PM's and devs involved
- Put features behind toggles
- Maintenance mode(s)
- Read-only mode





Preparation For the Worst

Embrace your fears:

- Think about what can go wrong
- Plan a response - write Runbooks!
- Practice

I don't always test in production...



It's Not All Bad

Sometimes you have to do this anyway:

- Upgrades
- New hardware
- New features mean new tables & queries

Try to use the same procedures every time.



What Could go Wrong?

- Host failures
- Network failures
- Provider failures
- Overwhelming usage





Best Practices - Server

Normalize maintenance & disaster recovery

- Replacing a replica
 - Same as building a new one
- Replacing a master
 - Same as promoting a new one
- Cloning a database
 - Same as restoring a backup



Best Practices - Application

- Give developers insight into DB health during:
 - Pushing code
 - Traffic surges





Learn From Mistakes

- Blameless Post-mortems
- Retrospectives



EURUSD - 1,35379 - 00:00:00 14 giu (EEST)

EURUSD (Bid), Ticks, # 300 / 300



Gold, spot - 1.276,820 - 23:00:00 13 giu (CEST)

Gold, spot (Bid), 1 minute, # 159 / 300, Logarithmic, Heikin Ashi



Advanced Monitoring and Investigation

- Monitoring & alerting is pretty much
 - Collecting data
 - Using it



Monitoring the Server

Examples of things to monitor on a DB server:

- Is MySQL running?
- ~~Load~~ % CPU Idle
- Disk usage
- Critical processes
- Critical jobs



Monitoring MySQL

- Auto-increment id limits
- Binary log growth
- Killed queries
- Replication delay, hierarchy
- `read_only`
- Temporary tables on disk
- Threads connected
- Threads running
- Fancy replication checks



Monitoring MySQL Usage

SHOW GLOBAL STATUS has counters & gauges

- Queries per second
- Threads running
- Threads connected
- Slow queries
- Open transactions



Monitoring MySQL Client

You can monitor these through our client:

- Error rate
 - Rates per type
- Log all warnings & errors





<https://www.flickr.com/photos/wocintechchat/25392405123/>

Alerting

Ways your monitors can send you info:

- ~~Email~~
- Chat
- Tickets
- Pages

Chatbots are so hot right now.



Send Good Alerts



Make your Alerts Actionable

Make sure the alert:

- Goes to the right team!
- Contains
 - A link to the runbook/guide
 - The first step to take



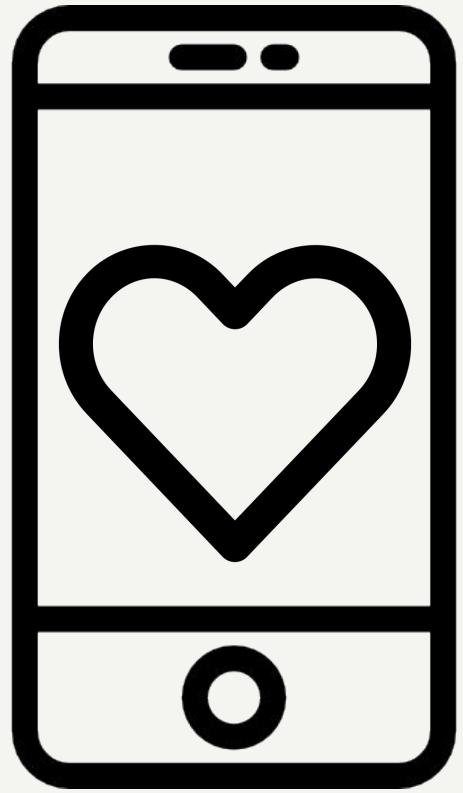
Iterate and Pay it Forward

Nothing is perfect the first time around

- Tune thresholds
- Keep track of alerts and response
- Handover to the next person down the line

If it happens consistently, automate the response!







Have a Runbook/Guide

- Start with the read some logs/part
- Iterate, refine, shorten



Tracking More Than Metrics

Cheapa\$\$ cron jobs capture this every minute:

- ps
- mysql -e "SHOW PROCESSLIST"
- mysql -e "SHOW ENGINE INNODB STATUS"

And this once a day:

- mysql -e "SHOW TABLE STATUS"
- mysqldump --no-data

What happened just before I was paged?



```
jsnyder@db:~$ cat ps-log.sh
```

```
#!/bin/bash

[ -d /var/log/ps ] || mkdir /var/log/ps

(
date -R
/bin/ps auxf
free
echo
) | bzip2 -c >> /var/log/ps/ps-`/bin/date +%Y%m%d-%H`.bz2
```

pt-stalk

Collect forensic data about MySQL when problems occur.

- Can run as a daemon
- Highly configurable
- Waits for a trigger condition
- Collects a ton of information





Fantastic Errors...

- Connection Errors
- Client Errors
- Server Errors
- Replication Errors

2013, 'Lost connection to MySQL server during query'



And Where to Find Them

- Your application's log
- MySQL's error_log
- SHOW SLAVE STATUS
- SHOW ENGINE INNODB STATUS

Just <insert search engine here> for it...



```
Traceback (most recent call last):
```

```
  File "/blahblah/blah/blah.py", line 357, in run
    return self.execute()
  File "/blah/blah.py", line 359, in execute
    return action(**match)
  File "/blah/blaah.py", line 172, in blaah_html
    return self.xtmpl('blaah.blaah',
env=self.blaah_env(blaah_form_obj))
  File "/lib/blah-blah/orm/session.py", line 470, in __exit__
    self.rollback()
  File "/lib/blah-blah/bleh.py", line 238, in commit
    self._delegate.commit()
```

```
OperationalError: (2013, 'Lost connection to MySQL  
server during query') None None
```

Debugging Lost Connections

- Try running the query itself
 - Does it return a huge number of rows? Sometimes?
 - Return a large amount of data?
 - Increase `max_allowed_packet`
- Does it go away if you increase `net_read_timeout`?
- Is the application server overloaded?
- It *might* be the network



MySQL Goes Away...

- Sometimes MySQL goes away...

OperationalError: (2006, 'MySQL server has gone away')





When an Error Contains a Query

A sensitive subject:

- You may store sensitive information in your DB
 - Email address
 - First & last names
 - Payment data
- You probably don't want it in your logs
- [pt-fingerprint](#)



Server Errors

You'll find these in the error_log

- log_error_verbosity = 1
- start/stop messages
- replication information
- crashes



InnoDB: Warning: a long semaphore wait:

--Thread 140375092094720 has waited at row0purge.cc line 770
for 241.00 seconds the semaphore:

S-lock on RW-latch at 0x1346820 '&dict_operation_lock'
a writer (thread id 140375041206016) has reserved it in mode
exclusive

number of readers 0, waiters flag 1, lock_word: 0

Last time read locked in file row0purge.cc line 770

Last time write locked in file

/dist/work/percona-server/storage/innobase/row/row0mysql.cc
line 3827

InnoDB: ##### Starts InnoDB Monitor for 30 secs to print
diagnostic info:

InnoDB: Pending preads 5, pwrites 0

MySQL Crashes

- MySQL will dump what info it can to the log
- `mysqld_safe` will attempt to restart MySQL
- Crashes are not normal



16:22:00 UTC - mysqld got signal 11 ;

This could be because you hit a bug. It is also possible that this binary or one of the libraries it was linked against blah blahblah blah blah blah blah.

```
key_buffer_size=268435456  
read_buffer_size=131072  
max_used_connections=8  
max_threads=41  
thread_count=2  
connection_count=2
```

It is possible that mysqld could use up to

$\text{key_buffer_size} + (\text{read_buffer_size} + \text{sort_buffer_size}) * \text{max_threads}$ = 309930 K bytes of memory

Hope that's ok; if not, decrease some variables in the equation.

Common Crash Causes

- OOM Killer
- Corrupted Data
- Hardware Failure
- Bug





When Things Are "Slow"

Lots of things can cause this:

- Individual statements can be slow
- Lots of queries
- Lots of slow-ish queries

Find out what resource is in contention.



Tuning an Individual Query or Table

Examples here are from our open dataset.

Not queries we use in real-life :)



```
# Time: 2017-10-27T19:38:53.671918Z
# User@Host: root[root] @ localhost [] Id:      13
# Query_time: 4.947666  Lock_time: 0.000416 Rows_sent: 5
Rows_examined: 610939
SET timestamp=1509305933;
select count(*) , r.stars
from review r
join business b on r.business_id = b.id
join category c on b.id = c.business_id
where c.category = 'Soul Food'
group by r.stars ;
```

```
mysql> select count(*), r.stars from review r join
business b on r.business_id = b.id join category c on
b.id = c.business_id where c.category = 'Soul Food'
group by r.stars ;
+-----+-----+
| count(*) | stars |
+-----+-----+
|      2052 |     1 |
|     1785 |     2 |
|     2356 |     3 |
|     5432 |     4 |
|    8779 |     5 |
+-----+-----+
5 rows in set (4.95 sec)
```

```
mysql> EXPLAIN select count(*), r.stars from review r join
business b on r.business_id = b.id join /*snip*/ ... \G
***** 1. row *****
      id: 1
select_type: SIMPLE
      table: category
    partitions: NULL
        type: ALL
possible_keys: fk_categories_business1_idx
          key: NULL
      key_len: NULL
          ref: NULL
        rows: 587969
  filtered: 10.00
    Extra: Using where; Using temporary; Using
filesort
```

```
***** 2. row *****
    id: 1
select_type: SIMPLE
    table: business
partitions: NULL
    type: eq_ref
possible_keys: PRIMARY
    key: PRIMARY
key_len: 68
    ref: yelp_db.c.business_id
rows: 1
filtered: 100.00
Extra: Using index
```

```
***** 3. row *****
      id: 1
select_type: SIMPLE
      table: review
partitions: NULL
      type: ref
possible_keys: fk_reviews_business1_idx
      key: fk_reviews_business1_idx
key_len: 68
      ref: yelp_db.c.business_id
     rows: 25
filtered: 100.00
     Extra: NULL
3 rows in set, 1 warning (0.00 sec)
```

```
mysql> pager fgrep rows
PAGER set to 'fgrep rows'
```

```
mysql> EXPLAIN select count(*), r.stars from review r join
business b on r.business_id = b.id join category c on b.id =
c.business_id where c.category = 'Soul Food' group by r.stars
\G
```

```
rows: 587969
```

```
rows: 1
```

```
rows: 25
```

```
3 rows in set, 1 warning (0.00 sec)
```

```
mysql> nopager
PAGER set to stdout

mysql> select 587969 * 1 * 25 ;
+-----+
| 587969 * 1 * 25 |
+-----+
|          14699225 |
+-----+
1 row in set (0.00 sec)
```

```
# From our slow query log, the numbers can be pretty off!
# Query_time: 4.947666  Lock_time: 0.000416 Rows_sent: 5
Rows_examined: 610939
```

```
mysql> alter table category add key (category) ;  
Query OK, 0 rows affected (1.53 sec)  
Records: 0  Duplicates: 0  Warnings: 0
```

```
mysql> alter table review drop key fk_reviews_business1_idx ,  
add key( business_id, stars) ;  
Query OK, 0 rows affected (38.23 sec)  
Records: 0  Duplicates: 0  Warnings: 0
```

```
mysql> alter table category add key (category) ;
```

```
Query OK, 0 rows affected (1.53 sec)
```

```
Records: 0  Duplicates: 0  Warnings: 0
```

```
***** 1. row *****
```

```
    id: 1
```

```
select_type: SIMPLE
```

```
    table: c
```

```
partitions: NULL
```

```
    type: ref
```

```
possible_keys: fk_categories_business1_idx,category
```

```
    key: category
```

```
key_len: 768
```

```
    ref: const
```

```
    rows: 235
```

```
filtered: 100.00
```

```
Extra: Using temporary; Using filesort
```

```
mysql> select count(*), r.stars from review r join business b  
on r.business_id = b.id join category c on b.id = c.business_id  
where c.category = 'Soul Food' group by r.stars ;
```

count(*)	stars
2052	1
1785	2
2356	3
5432	4
8779	5

```
5 rows in set (0.03 sec)
```

Changing the Schema IRL

- Long ALTER TABLE's can hurt performance
- [pt-online-schema-change](#)
- [gh-ost](#)



And Just a Bit More on Indexes

There's a [a great manual page How MySQL Uses Indexes](#)

PERFORMANCE_SCHEMA can give you stats on index usage.

Keep in mind:

- *Adding indexes slows down writes a bit*
- *Cardinality*
- *Leftmost-prefixing*







Lots of Queries at Once

Find them using SHOW PROCESSLIST

Or SELECT ... FROM information_schema.processlist

You can use cheap bash commands to sort & find culprits.



```
mysql -uroot -e "show full processlist" | awk '{print $5}' |  
sort | uniq -c | sort -gr | head
```

```
jnyder@db:~$ mysql -uroot -e "show full processlist" | awk  
'{print $5}' | cut -c 1-70 | sort | uniq -c | sort -gr | head  
521 Sleep  
109 select count(*), r.stars from review r j  
3 Query  
2 Connect  
1 Command
```

```
# kill them with pt-kill. Use --print to test first!  
pt-kill --config /etc/mysql/emergency-kill.conf --kill  
--match-info 'your query' --daemonize --run-time=60
```

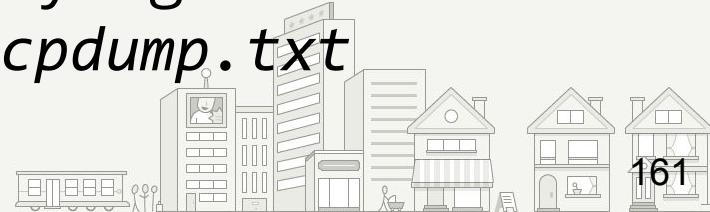
Finding the Next Thing To Optimize

You can:

- Some slow queries run infrequently
- Some slowish queries run all the time
- Visualization tools like Anemometer, VividCortex

Take a tcpdump & analyze with pt-query-digest:

[*pt-query-digest*](#) --type tcpdump tcpdump.txt





Tuning MySQL/InnoDB

- I/O
- Memory
- Tables/Transactions



I/O is Always the Slowest

InnoDB has a ton of knobs and dials

- [Innodb io capacity](#)
- Optimize your filesystem
- Add more memory, larger/more buffer pools
- Tinker with [innodb flush method](#)

*There is a whole manual page dedicated to optimizing
[InnoDB disk I/O.](#)*



How InnoDB Uses Memory

Think of the InnoDB buffer pool as InnoDB's cache

- Increase the amount of memory to reduce I/O

Rough math:

$$\begin{aligned} & \text{innodb_buffer_pool_instances} * \text{innodb_log_buffer_size} \\ & + (\text{read_buffer_size} + \text{sort_buffer_size}) * \text{max_threads} \end{aligned}$$

[MySQL Memory Calculator](#) looks out of date, but can give you a pretty good idea...



On Transactions

To view running transactions:

- SHOW ENGINE INNODB STATUS
- SELECT * FROM information_schema.INNODB_TRX \G



```
mysql> SHOW ENGINE INNODB STATUS ;
```

```
...
```

```
-----
```

```
TRANSACTIONS
```

```
-----
```

```
Trx id counter 50806393604
```

```
Purge done for trx's n:o < 50806389839 undo n:o < 0 state:  
running but idle
```

```
History list length 527
```

```
LIST OF TRANSACTIONS FOR EACH SESSION:
```

```
...
```

...

---TRANSACTION 50806392298, **ACTIVE 12** sec starting index read
mysql tables in use 1, locked 1

**314 lock struct(s), heap size 46632, 358 row lock(s), undo log
entries 353**

MySQL thread id 23649033, OS thread handle 0x7f6bd31c7700,
query id 32707038971 10.10.10.40 application_user updating
UPDATE business SET ...

Trx read view will not see trx with id >= 50806392299, sees <
50806389830

Improving Transaction Efficiency

Reduce lock contention by:

- Decreasing the number of locks your transactions use
- The total time of your transactions
- Making writes faster

See [InnoDB Locking Explained with Stick Figures](#)

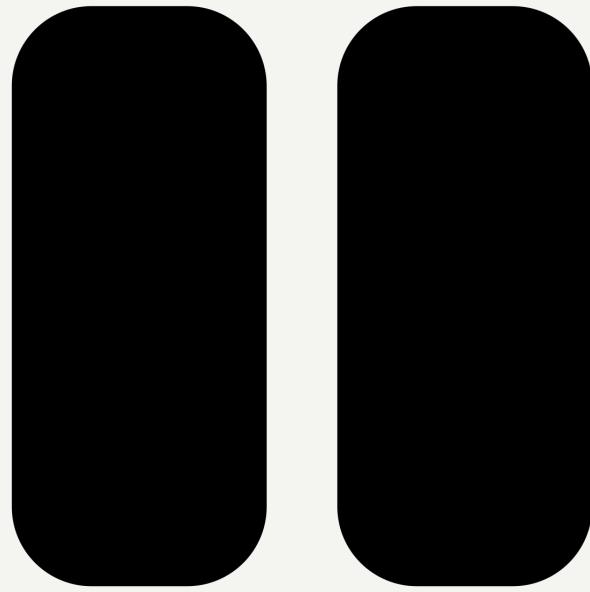


Other Ways to Look at Transactions

- Performance_schema
 - table_lock_waits_summary_by_table



Let's Hit Pause For a Moment





Replication and Scaling Out

Ways to use replication:

- Take load off the master
- Send reads to the replicas
- Having a separate reporting database
- Delayed replica with [pt-slave-delay](#)



How MySQL Replication Works Again

- Master writes binary logs
- Each replica has 2 threads:
 - IO_THREAD: reads from master, writes relay logs
 - SQL_THREAD: runs SQL from relay logs
- Serial per database with parallel replication

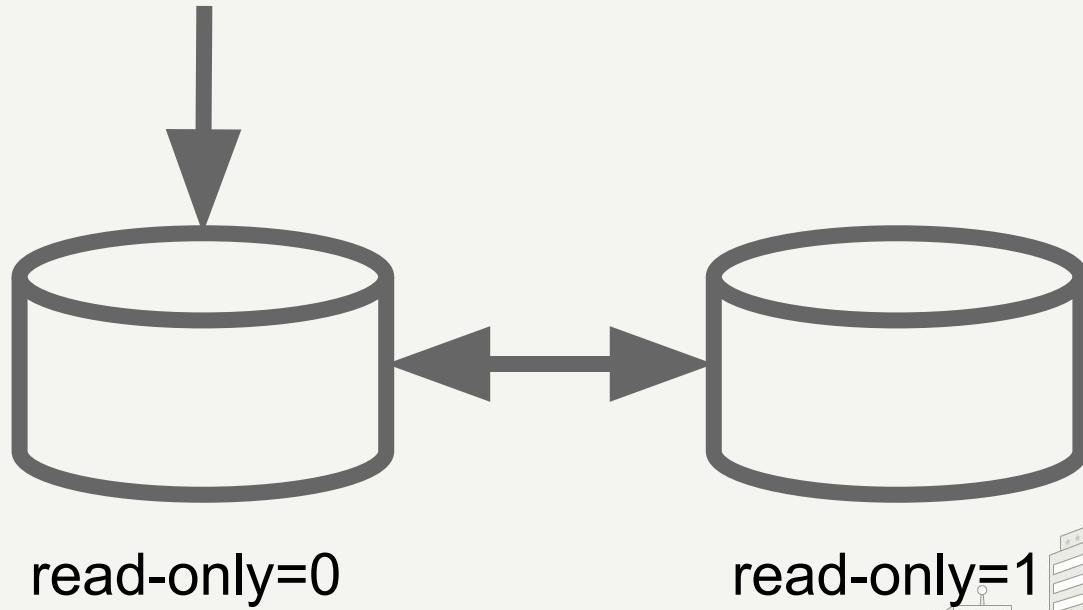


Some Replication Strategies

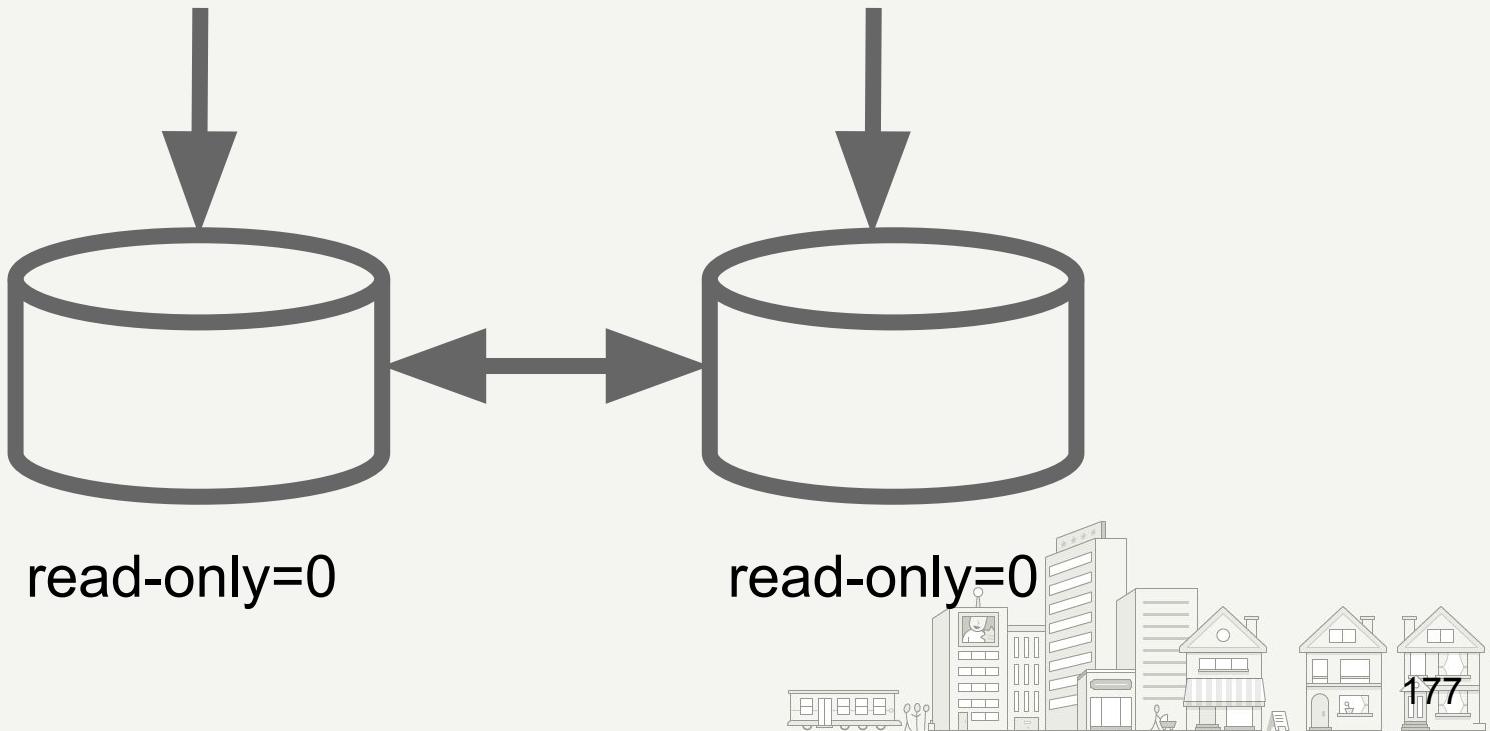
- master-master
 - One master "live", or
 - `auto_increment_increment`, offset
- master, intermediate master
- master(s) with many replicas



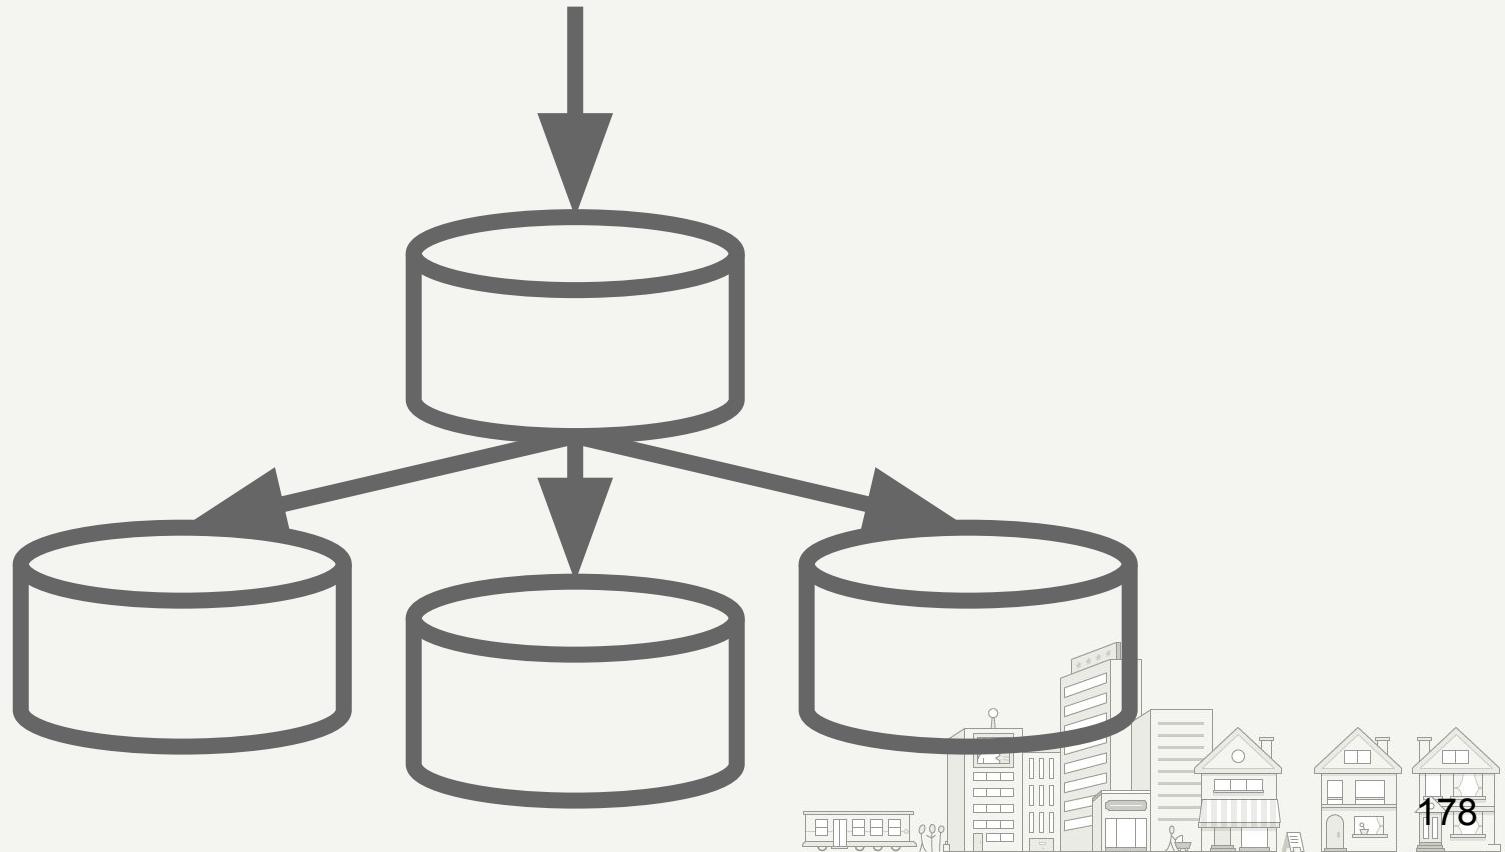
Master-Master - Active-Standby



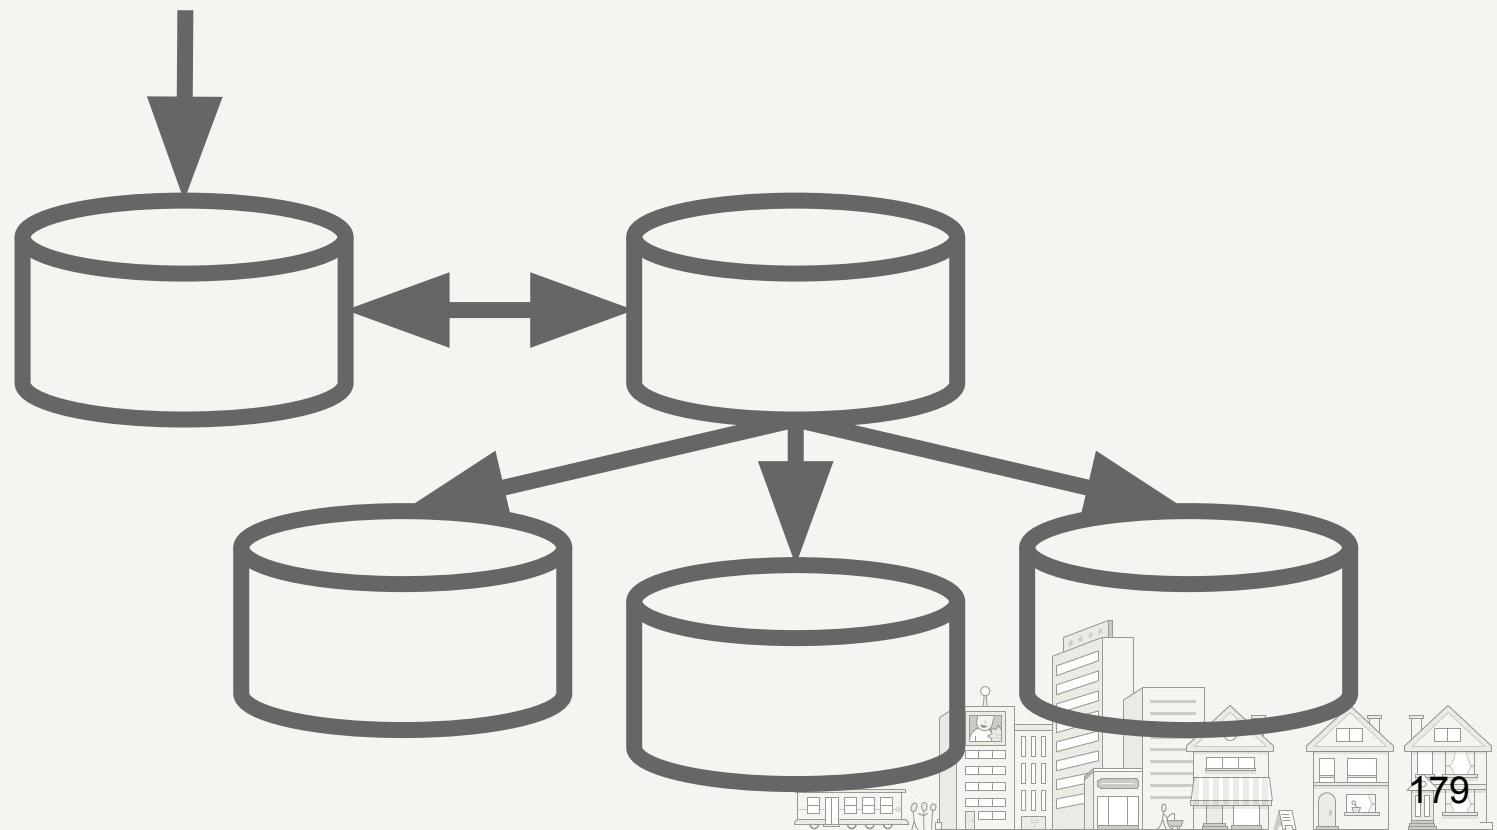
Master-Master - Active-Active



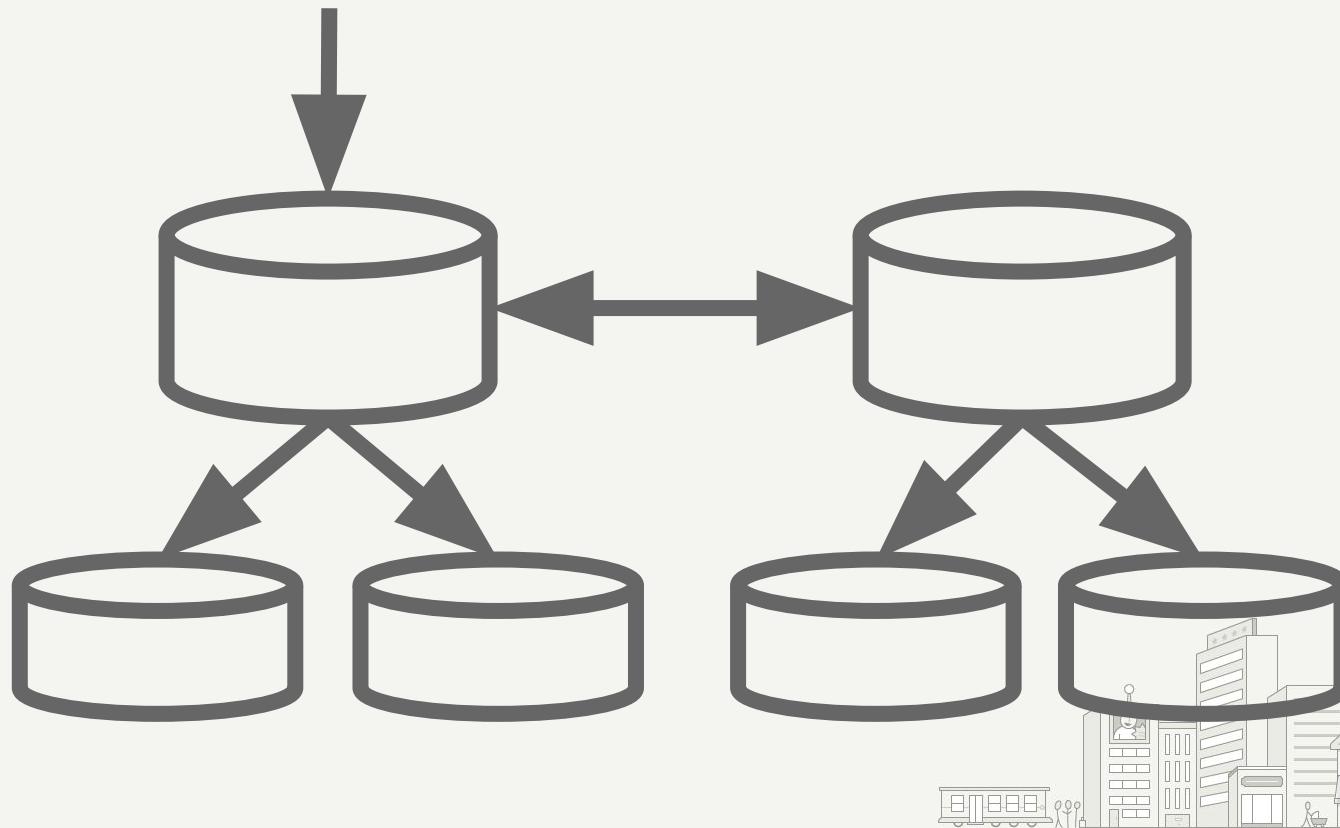
One Master, Many Replicas



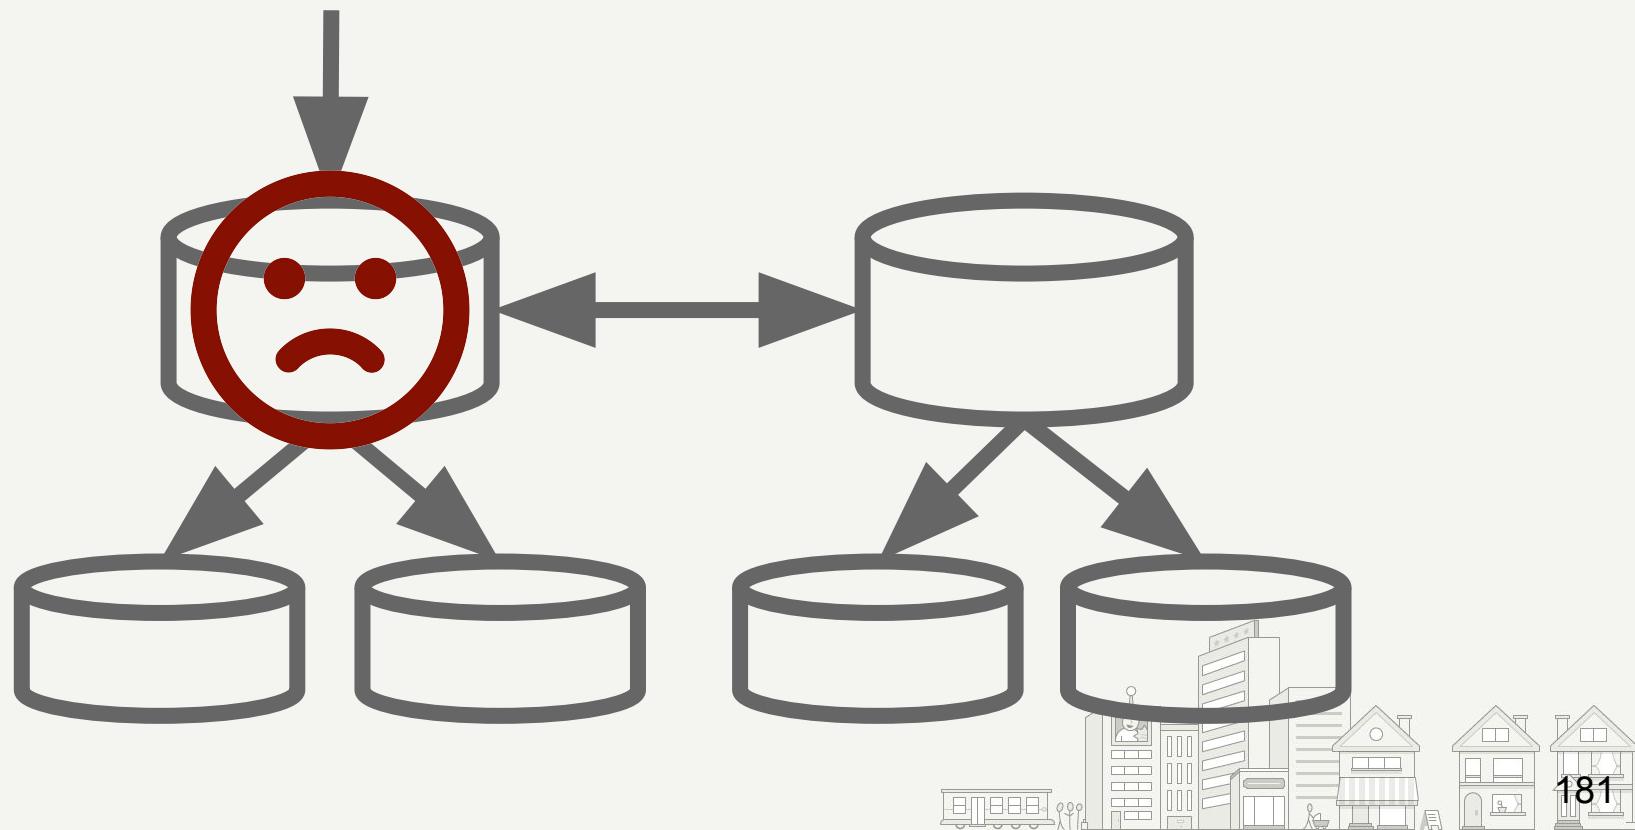
master, intermediate master



masters with their own replicas



Think About What You Want to Do



Tools For Replication

- Enable GTID - easy reparenting
- [orchestrator](#) - GUI and tooling
- [pt-table-checksum](#) - check consistency
- [ProxySQL](#) - can help you split your read-write traffic





When in Doubt, Throw it Out

- Don't write to sketchy databases
- Don't let your application write to replicas
- use `read_only=1` to prevent writes to replicas
- `SHOW MASTER STATUS, SHOW SLAVE STATUS`

Remember, a host that goes away may wander back...



Sample Master Promotion Plan

1. Set old master to `read_only=1`
2. `CHANGE MASTER TO` if necessary
3. Configure your application to write to new master
4. Set new master to `read_only=0`

Never allow writes to a replica unless you know that's what you want!





Speaking of Lag...

Common causes:

- IO_Thread can't keep up
 - Local load or slow/lossy network
- SQL_Thread can't keep up
 - Local load, long statements or transactions
 - mysqlbinlog to look at transactions



When You're Really Stuck

You can always phone a friend:

- MySQL has a big, friendly online community
- Oracle MySQL, Percona, MariaDB offer support

Get the most out of your interactions with experts by providing lots of context and what you've already tried.



**Questions
Answers**

Let's Talk About It

We covered:

- Running in Production
- Being Prepared
- Advanced Monitoring and Investigation
- Fantastic Errors and Where to Find Them
- When Things Are "Slow"
- Advanced Tuning
- Replication and Scaling Out





We're Hiring!

www.yelp.com/careers/



[fb.com/YelpEngineers](https://www.facebook.com/YelpEngineers)



[@YelpEngineering](https://twitter.com/YelpEngineering)



engineeringblog.yelp.com



github.com/yelp

Appendix - Links to Resources

- [MySQL Manual](#)
- [MySQL Utilities](#)
- [Percona Toolkit](#)
- [Orchestrator: MySQL Replication Topology Manager](#)
- [ProxySQL: High Performance MySQL Proxy](#)
- [mysql-sys - tools for using performance_schema](#)
- [Anemometer - Slow Query Monitor](#)



Appendix - Books

- *Database Reliability Engineering* by Laine Campbell, Charity Majors
- *High Performance MySQL* by Baron Schwartz, Peter Zaitsev & Vadim Tkachenko



Appendix - Blog Posts

- [Database Defense in Depth by @geodbz](#)
- [GrossQueryChecker: Don't get surprised by query performance in production! by @jcsuperstar](#)
- [InnoDB Locking Explained with Stick Figures by @billkarwin](#)
- [Blameless PostMortems and a Just Culture by @allspaw](#)
- Literally any [Percona Blog Post by @svetsmirnova](#)



Appendix - Image Sources

- [Pexels - Images licensed under CC0](#)
- Icons from [The Noun Project](#)
- [#WOCInTechChat](#)
- [The Jopwell Collection](#)



Appendix - Yelp Dataset

- [Yelp Dataset](#)

Other datasets:

- [MySQL's Employees Sample Database](#)
- [Percona's list of Sample datasets](#)

